



# **IT DATA SECURITY LAB FILE**

**Name- Dhairya Jain**  
**Sap ID- 500105432**  
**Batch- CSF-B4**

## EXPERIMENT-11

### Part A- GPU Parallel Programming

#### Example Using CuPy

- Install CuPy

```
D:\data sec>pip install cupy-cuda11x
Requirement already satisfied: cupy-cuda11x in c:\python312\lib\site-packages (13.3.0)
Requirement already satisfied: numpy<2.3,>=1.22 in c:\python312\lib\site-packages (from cupy-cuda11x) (2.0.2)
Requirement already satisfied: fastrlock>=0.5 in c:\python312\lib\site-packages (from cupy-cuda11x) (0.8.2)

[notice] A new release of pip is available: 24.0 -> 24.3.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

- GPU Parallel Programming Example Using CuPy  
Code-

```
exp11_1.py X exp11_2.py hello.py
exp11_1.py > ...
1  import cupy as cp
2  import numpy as np
3  import time
4  # Define the size of the arrays
5  n = 10000000
6  # Create two large arrays on the GPU
7  a_gpu = cp.random.rand(n)
8  b_gpu = cp.random.rand(n)
9  # Perform the addition on the GPU
10 start_gpu = time.time()
11 c_gpu = a_gpu + b_gpu
12 end_gpu = time.time()
13 # Transfer the result back to the CPU (if needed)
14 c_cpu = cp.asnumpy(c_gpu)
15 print("GPU addition took:", end_gpu - start_gpu, "seconds")
16 # For comparison, perform the same operation on the CPU using NumPy
17 a_cpu = np.random.rand(n)
18 b_cpu = np.random.rand(n)
19 start_cpu = time.time()
20 c_cpu = a_cpu + b_cpu
21 end_cpu = time.time()
22 print("CPU addition took:", end_cpu - start_cpu, "seconds")
```

#### Output-

```
PS D:\data sec> python -u "d:\data sec\exp11_1.py"
GPU addition took: 0.006550788879394531 seconds
CPU addition took: 0.027991056442260742 seconds
PS D:\data sec>
```

- Explanation:
  - Creating Arrays: Two large random arrays are created directly on the GPU using CuPy.
  - Parallel Addition: The addition operation is performed on the GPU, leveraging its parallel processing capabilities.
  - Performance Comparison: The script also performs the same operation on the CPU using NumPy to compare performance.

## Example Using Numba

- Install Numba and CUDA Toolkit

```
D:\data sec>pip install numba
Requirement already satisfied: numba in c:\python312\lib\site-packages (0.60.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in c:\python312\lib\site-packages (from numba) (0.43.0)
Requirement already satisfied: numpy<2.1,>=1.22 in c:\python312\lib\site-packages (from numba) (2.0.0)

[notice] A new release of pip is available: 24.0 -> 24.3.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

- GPU Parallel Programming Example Using Numba Code-

```
exp11_1.py  exp11_2.py x  hello.py
exp11_2.py > vector_add
1  import numpy as np
2  from numba import cuda
3  import time
4  # Define the size of the arrays
5  n = 10000000
6  # Initialize arrays on the CPU
7  a_cpu = np.random.rand(n)
8  b_cpu = np.random.rand(n)
9  c_cpu = np.zeros(n)
10 # Define a GPU kernel for vector addition
11 @cuda.jit
12 def vector_add(a, b, c):
13     idx = cuda.grid(1)
14     if idx < a.size:
15         c[idx] = a[idx] + b[idx]
16 # Allocate memory on the GPU
17 a_gpu = cuda.to_device(a_cpu)
18 b_gpu = cuda.to_device(b_cpu)
19 c_gpu = cuda.device_array_like(a_cpu)
20 # Define the number of threads and blocks
21 threads_per_block = 1024
22 blocks_per_grid = (a_cpu.size + (threads_per_block - 1)) // threads_per_block
23 # Perform the addition on the GPU
24 start_gpu = time.time()
25 vector_add[blocks_per_grid, threads_per_block](a_gpu, b_gpu, c_gpu)
26 cuda.synchronize()
27 end_gpu = time.time()
28 # Transfer the result back to the CPU
29 c_cpu = c_gpu.copy_to_host()
30 print("GPU addition took:", end_gpu - start_gpu, "seconds")
31 # For comparison, perform the same operation on the CPU
32 start_cpu = time.time()
33 c_cpu = a_cpu + b_cpu
34 end_cpu = time.time()
35 print("CPU addition took:", end_cpu - start_cpu, "seconds")
```

Output-

```
PS D:\data sec> python -u "d:\data sec\exp11_2.py"  
GPU addition took: 0.19643712043762207 seconds  
CPU addition took: 0.027849912643432617 seconds  
PS D:\data sec>
```

- Explanation:
  - Defining a GPU Kernel: The `vector_add` function is decorated with `@cuda.jit`, indicating it should be compiled to run on the GPU. The `cuda.grid(1)` function provides the unique index of each thread in the 1D grid.
  - Memory Management: Arrays are transferred from the CPU to the GPU and vice versa using `cuda.to_device()` and `cuda.device_array_like()`.
  - Kernel Execution: The kernel is launched with a specific number of blocks and threads, allowing it to run in parallel on the GPU.
  - Synchronization: `cuda.synchronize()` ensures that the GPU has completed all its work before measuring the execution time.